

Speaker 3 — The Three Doors

TEAM RUBIXCUBE · SIH 2026 · PS 26097 · Idea video — your section only

⚠ The one rule for this video

We did no field work for this problem statement. Nobody says otherwise on camera.

Never say “we surveyed”, “we visited”, “we interviewed” or “people told us”. Our credibility comes from primary-source documents we read in full — a CAG performance audit, the NSQF gazette, PM-AJAY’s own guidelines and the national qualification register. Attribute to *them*, never to us.

If anyone ad-libs a field-work claim, the take is dead — re-shoot it.

Speaker 3 | 1:22 – 2:34 | ~180 words

This is the demo section, and it is the section that changed most. It used to walk only the phone call. The deck draws three bridges, so this now walks all three — the call in full, then WhatsApp and the app as two short beats that prove the same engine is behind them. The IVR call is real and filmed; WhatsApp and the app are screen-recorded off the real build; the interview internals are an animated walkthrough built in post.

Bullet points (what they must land)

- **Do not re-list features. Walk one person through one call, start to finish — then show the other two doors happening to other people.** Name her. Tag her **ILLUSTRATIVE** on screen the first time the name appears.
- **Budget discipline:** the call is ~48 seconds, WhatsApp ~10, the app ~7, the close ~7. The two extra doors are *beats*, not tours. If you give them 20 seconds each you lose the near-miss, which is the best thing in the section.

Door 1 — the phone call (the long one). In this exact order:

1. She gives a **missed call** from a feature phone. **No smartphone, no data, no app. The system rings her back — so the call costs her nothing.** This is the flow that is actually built, and for a poverty-targeted scheme it is a *better* line than “toll-free”.
2. The system knows her from the phone number — nothing to read out, nothing to spell.
3. It asks, in her language. She answers in hers.
4. **It reads her answer back and waits for a yes.** Say “**nothing is counted as her answer until she confirms**” — precise and true. Do **not** say “it saves nothing until then”.
5. Seven questions, one at a time, one per turn.
6. **The eligibility gate fires** — it filters the national register down to what she can actually enrol in, using her schooling *and* her years of work.
7. **The near-miss** — for the trade she wanted but can't enter yet, it names the exact gap. This is the problem statement's “skill gaps requiring intervention”, answered literally. **Name the gap, never a cure.** There is no bridge-course recommendation in the build.
8. She gets the answer spoken back: the trade, how long the training runs, and the **₹50,000 asset grant** the scheme entitles her to. (*All three are in the real output. The nearest centre is not — do not say it.*)
9. **The call can drop and she can call back and continue** — on a four-digit PIN. Shared handsets and bad rural lines are the normal case, not the edge case.

Door 2 — the WhatsApp voice note (~10s). A different person, not Sunita — the point is that the channel is a *choice*, not a fallback. She holds the mic, sends a voice note, and **one question comes back as audio**. Two details worth the words, because they are design decisions and a judge will hear them as such:

- **One question per message.** We costed batching and rejected it: batching turns WhatsApp into a form read aloud, which is the thing this whole product exists to avoid.
- **She answers between chores.** The session waits. It is the asynchronous door — that is what it is *for*, and it is the reason to have it alongside a live call.
- *(Not on camera, but know it: resume on this channel is forward-only and never reads prior answers back, because a chat log lives on a handset that is frequently shared.)*

Door 3 — the kiosk app (~7s). The CSC operator — the Village Level Entrepreneur — runs the **same interview** on the app, for someone who has no phone of her own at all. The one thing to say about it: **it keeps going when the link drops.** Answers are held on the device and sync later. Say **"it keeps working when the link drops", never "it works offline"** — the interview and the record are offline, the microphone is not yet. That distinction is in the three-doors rule at the top and it is the only trap in this section.

The close. "Three doors. One engine. One record." This is the line that makes the section about architecture rather than about features, and it is exactly what the Technical Approach slide draws: three bridges converging on one Server Processes box. Then the old closer, unchanged: **"Nobody read anything. Nobody typed anything."**

- **The honesty beat still lives here, spoken, not in a caption:** the register we recommend from is the official one, and where the official record has no value, we store nothing rather than guess.

SAY THIS 180 words ≈ 72s

"Sunita gives a missed call from a feature phone. No smartphone. No data. No app. The system rings her back — the call costs her nothing.

It knows her from the number, asks in her language, and repeats what it heard until she says yes. Nothing counts as her answer until she confirms.

Seven questions. One at a time.

Then it filters the national register down to what she can actually enrol in — her schooling, and her years at a loom. For the trade she can't enter yet, it names exactly what she is short of.

Then it speaks the answer back — the trade, the training length, and the fifty-thousand rupee asset grant.

Her line drops; she calls back, keys a PIN, and picks up where she left off.

Her neighbour never calls at all. She holds the mic in WhatsApp — one question back as audio, one at a time, answered between chores.

At the CSC, the operator runs the same interview in the app; it keeps going when the link drops.

Three doors. One engine. One record.

Nobody read anything. Nobody typed anything."

Delivery note: Conversational. This is a story, not a spec. The WhatsApp and app lines are said **faster and flatter** than the call — they are evidence, not a second story, and the pace change is what tells the viewer that. Then full stops between "Three doors. One engine. One record." and again between the last two sentences.

⚠️ **All three doors are present tense. No watermark on any of them.** The build state and the two seams are in the three-doors table at the top of this document — read it before the shoot, because the person who says this on camera is the person a judge will ask.

⚠️ **Sunita is illustrative and must be labelled on screen.** Because Section 5 is built on real audited evidence, a judge will assume any named person is real. One small **ILLUSTRATIVE** tag the first time her name appears removes the risk entirely. The WhatsApp and app people are unnamed on purpose — one illustrative name in the video is enough.

⚠️ **The IVR call is real — lead with it and let the audience see it is real.** Film the actual handset, the actual ringing, the actual answer. Everything after that is screen-record or animated walkthrough. The cut between them must be obvious, not disguised.

⚠️ **Do not show a district officer screen in this section.** Section 4 handles the officer half; keep the two visually separate so each lands on its own.

📌 **Five phrases that are false against the current build. Do not let them back in:** 1. **"toll-free"** — the repo costed it and chose a normal number. Nothing is provisioned. The implemented flow is missed-call → the system rings *her* back. 2. **"the bridge course that closes it"** — no bridge-course logic exists anywhere in the engine. The output names the gap *she* must close; it does not prescribe a cure. 3. **"the nearest centre"** — not in the spoken output. Title, level, duration, one reason, the financial-literacy line and the ₹50,000 grant are. 4. **"only then does it save anything"** — the pending value is written to disk every turn, and the app path persists unconfirmed answers *with transcripts* by design. The honest claim is about what **counts**, not what is **stored**. 5. **"it works offline" / "in aeroplane mode", said about speech.** The app's microphone goes through Android's own recogniser, which usually needs a network hop. Vosk on-device is the missing piece. "It keeps going when the link drops" is true of the interview and the record, and it is the sentence in the script for a reason.

⚠️ **"Seven questions" is true of the mandate, not of the screen.** A filmed call shows language select, consent, PIN, then **nine** field prompts (the seven PS fields plus district and years-of-experience). Say "the seven things the problem statement asks for" — already the wording in Section 2 — and don't invite a count. If the edit shows a progress rail, don't label it "1 of 7" over a nine-step call.

⚠️ **If the filmed take uses the keypad, the read-back chip will not fire.** Keypad answers skip confirmation by design because they are exact. **Shoot a spoken-answer turn** for the read-back beat, or the voiceover and the footage will disagree on screen.

🎬 **ASSET CUE (Utsav):** Heaviest asset section, and it now carries three sources of footage — keep them visually distinct so the viewer can tell filmed from recorded from animated.

(a) The call — real footage, ~30s of screen time. Open on the real filmed call: a real feature phone, a real hand, a real **missed call**, then the phone ringing *back*, then the system answering. Hold it long enough to register as unfaked (3–4s). Then move to the animated walkthrough for the interview turns. Beats to capture: the read-back-and-confirm chip, the seven-question progress rail, the eligibility filter greying out ineligible qualifications, the **NEAR MISS** card naming the exact gap, and the call dropping and resuming on the same state.

(b) WhatsApp — screen-record, ~10s. A real handset, real thread. The shot is: thumb on the mic, waveform, send — then an **audio reply bubble arriving and playing**, with the waveform moving. That last part is the whole point; a text bubble on screen while the voiceover says "voice note" undoes the section. One question visible per message, so the "one question per message" design decision is legible without being narrated.

(c) The app — screen-record, ~7s. The kiosk/mobile build: the interview screen with the mic active, an answer landing as a filled field, and **the offline indicator** as the link drops while the interview carries on. Do **not** put an aeroplane-mode toggle on screen next to a live microphone — that is the one frame that would contradict the build.

(d) The close. Three bridge icons — handset, WhatsApp waveform, kiosk — converging into one box, lifted from the Technical Approach slide so it reads as the same object. Speaker 3 is off-screen or a corner inset for the middle 50 seconds.

On the eligibility-filter animation — do not label the starting pile 2,814. That is the full NQR export. The gate never sees it: filtering to Level ≤ 4 leaves 1,838, and dropping rows already past **valid_till** leaves **1,198 live qualifications**

as of 2026-09-28. Label it 1,198 live qualifications and the narrowing is still the clearest visual in the video — and it is a number that survives someone checking it.