

Speaker 5 — The Evidence

TEAM RUBIXCUBE · SIH 2026 · PS 26097 · Idea video — your section only

⚠ The one rule for this video

We did no field work for this problem statement. Nobody says otherwise on camera.

Never say “we surveyed”, “we visited”, “we interviewed” or “people told us”. Our credibility comes from primary-source documents we read in full — a CAG performance audit, the NSQF gazette, PM-AJAY’s own guidelines and the national qualification register. Attribute to *them*, never to us.

If anyone ad-libs a field-work claim, the take is dead — re-shoot it.

Speaker 5 | 3:28 – 4:00 | ~86 words

Replaces the previous pack's "Ground Work" section. Read the rule at the top of this document before writing anything else into it.

Bullet points

- **The pivot: we did not invent this problem, and we did not survey it either. We read it.** The audit, the gazette, the guidelines, the qualification register — in full.
- Name the sources out loud. They are the credential: **a CAG performance audit, the NSQF gazette, the PM-AJAY guidelines, and the national qualification register.**
- **The honesty line, and it is the best line in the video:** we do not claim to have solved dialects. Models for these languages **do now exist** — and they are still wrong roughly **three words in ten** (Bhojpuri 27.8 WER, Chhattisgarhi 27.4, Magahi 30.4) and **four in ten** for Rajasthani (41.8), on ARTPARK-IISc's SraVaani-1.0, the best published numbers available. Together those four are **~10.5 crore speakers**, all non-scheduled. So we didn't wait for the models to get better. We built so that a **wrong transcript still produces a right field**, and **we publish our error rate next to our accuracy.**
- The close: return to the poster from Section 1. **This is the instrument.**
- **Final line is the team name.** End on the phone, not the team's faces.

SAY THIS 84 words ≈ 34s

"We didn't invent this problem. We read it — in a CAG audit, the NSQF gazette, PM-AJAY's own guidelines, and the national qualification register.

And one thing we will not claim: we have not solved dialects. The best published model for Bhojpuri still gets nearly three words in ten wrong. For Rajasthani, four in ten. So we built for the transcript being wrong — and we publish our error rate next to our accuracy.

The scheme asks that a beneficiary's interest be assessed. For twenty years the instrument for that was an advertisement.

This is the instrument.

Team RubixCube."

Delivery note: Quietest section. No energy push. "This is the instrument" gets a full stop before and after it.

 **This section contains zero field-work claims. Check it again before the take.** The previous pack's Section 5 was built on real interviews; this one is built on documents. If anyone ad-libs "we spoke to people", the take is dead — re-shoot it.

 **NEVER say "there is no speech model for Bhojpuri / Magahi / Chhattisgarhi."** An earlier draft of this script said exactly that and it is **false**. ARTPARK-IISc's SraVaani-1.0 publishes WER for all four. **Bhashini's own ASR list includes Bhojpuri and Chhattisgarhi** — and MoSJE funds Bhashini. That sentence, in front of this jury, would end the pitch. The claim is about **error rate**, never about **existence**.

 **The dialect admission is still a scoring line, and now a better one.** Every other team will claim multilingual support. We are the team that knows the 2026 numbers, quotes them against ourselves, and designed around them. "Models exist and they are still wrong three words in ten" is a more expert sentence than "no model exists" ever was.

 **"We publish our error rate" is true, but our published figures are gTTS-synthesised Hindi over a simulated line — not dialect audio.** 05-measurements.md says so itself. Do not let the voiceover imply we measured Bhojpuri. If pressed: *we publish the pipeline measurement we have, and we name what it does not yet cover.*

 **ASSET CUE (Utsav):** No people. Document montage — the CAG report cover, the gazette page, the guidelines page 26 calendar, the NQR export. Then the WER-vs-extraction-accuracy chart. **Check the real spread before you design this chart** — it is thinner than "one line bad, one line high". The measured worst case is 36% WER against 68% answers-understood; at noise10 it is 25% against 82%. Plot it honestly with both axes labelled; a ~40-point gap is still a clear story and an exaggerated one is the kind of thing this whole section exists to avoid. Last frame: the feature phone from Section 3, ringing. Titles over black: TEAM RUBIXCUBE – SIH 2026 – PS 26097 .

PART B — EDITOR SHEET (UTSAV)

B.1 — Timeline at a glance

#	Section	In-Out	Dur	Speaker	Visual mode	Face on cam?
1	The Problem	0:00–0:38	38s	S1	Document open + stat cards	First & last line only
2	What We Built	0:38–1:22	44s	S2	Three door-icons + eligibility filter	~50%
3	The Three Doors	1:22–2:34	72s	S3	Real filmed IVR call → animated walkthrough → WhatsApp screen-record → app screen-record	Corner inset / off for middle 50s
4	The Government Half	2:34–3:28	54s	S4	Big-number typography + real district & State console screen-record	Mostly ON camera
5	The Evidence	3:28–4:00	32s	S5	Document montage + WER chart	Opening line, then off

What changed from the previous cut, and why: Section 3 grew from 60s to 72s to carry the two extra doors the deck draws, paid for out of S1 (–2s), S2 (–1s), S4 (–6s) and S5 (–3s). The whole of that 12s goes to WhatsApp (~10s) and the app (~7s) minus the tightening inside the call itself. **Do not fund the two new beats by cutting the near-miss or the calendar** — those are the two things in this video other teams will not have.

Hard rule: never more than 8 seconds of a talking head with nothing else on screen. Never more than 12 seconds without a human face or a human voice doing something.

B.2 — Asset checklist

✓ **Have (real, film it)** - [] **The IVR call, filmed on a real feature phone.** Real hand, real **missed call**, then the phone ringing *back*, then the system answering. This is the highest-value asset in the entire video — everything else is animation, and this is the one shot that proves the thing exists. **Shoot the callback, not a dial-in** — that is the built flow, and the phone ringing by itself is the better image anyway. - [] **Shoot at least one spoken-answer turn**, not keypad. Keypad answers skip the read-back by design, and the read-back chip is a scripted beat in Section 3. - [] Time the response gap on the take. It decides whether Section 2 may say "under two seconds." - [] **Screen-record the WhatsApp thread on a real handset** — thumb on the mic, waveform, send, then **an audio reply bubble arriving and playing**. The playing waveform is the shot; a text bubble under a "voice note" voiceover kills the beat. One question visible per message. - [] **Screen-record the app** — the interview screen with the mic active, an answer landing in a field, and the **offline indicator appearing while the interview carries on**. No aeroplane-mode toggle on screen next to a live microphone. - [] **Screen-record /officer , /state and /outcomes** — all three are real, so all three belong here. `#/demo` seeds a district's worth of rows and drops straight into the console, which is the fastest way to record a console that has data in it. - [] Clean audio of one full turn from the demo call, in Hindi — question, answer, read-back, confirm. Subtitle it, never dub it. **Call it "the demo call", never "the interview"** in any caption or voiceover; "interview audio" is exactly the phrase a judge mishears as field work.

🔧 **Need to make (animated walkthrough, built in post)** - [] Interview turn animation: seven-question rail, read-back chip, confirm - [] **Eligibility filter animation** — the register narrowing, ineligible entries greying out. This is the single clearest visual argument in the video. Give it real screen time. - [] **NEAR MISS – exact gap** card - [] Call-drop-and-resume beat - [] District demand table building row by row - [] The four-date statutory calendar as a timeline - [] **Screen-record the real consoles** — **/officer 's five tabs, /state , /outcomes . No watermark; they are built.** Only the callback graphic carries a `NEXT` tag. Leave `/state` 's one-district honesty banner in shot. - [] **Three-bridges-into-one-box graphic** for the Section 3 close, lifted from the Technical Approach slide — handset, WhatsApp waveform, kiosk, converging on the engine. - [] PM-AJAY portal, if shown, tagged `EXISTING GOVERNMENT SYSTEM` - [] **WER-vs-extraction-accuracy chart** — Section 5 is built around it - [] Stat cards: `56.14 LAKH · 41% / 70% · 40% IN 10 ROLES · 90.35% → ONE ROLE · 3.5–4×`

📄 **Document footage (Section 1 and Section 5)** - [] CAG Report No. 20 of 2025 — cover and Table 2.1(a) - [] PM-AJAY Guidelines — the interest-assessment clause, and the page 26 calendar table - [] NSQF Gazette June 2023 — the entry-requirements table - [] NQR export — the 2,814-row register

All four are already in `docs/references/` or fetchable. Screen-record scrolling them; a real PDF scrolling reads as real research in a way a bullet list never does.

🤖 **AI-generated clips — maximum 2 in the whole video** 1. A woman on a feature phone in a rural setting — Section 3 establishing shot only 2. Optional: a training centre / workshop — Section 4

Rule: AI clips must never sit adjacent to the real filmed call. A synthetic shot two frames after the real handset kills the credibility of the real one.

B.3 — Audio

- One bed track, no vocals. Restrained.
- **Music drops to near-silence under the real IVR call audio.** That moment should sound like a phone call, not like a promo.
- Music also ducks under "Did you get work?" in Section 4 and under the entire Section 5 close.
- SFX budget, total: **one** dial tone, **one** soft confirm chime on the read-back, **one** transition on the calendar. Three. Anything more and it becomes a startup ad.

- Level lift of ~2dB going into Section 4.

B.4 — Graphics rule

Everything on screen must be liftable from the deck — same palette, same icons, same fonts as the Technical Approach slide. When a judge sees the flowchart after the video, it must read as the same object.

Specifically: use the deck's own three bridge icons — the Exotel/phone bridge, the WhatsApp bridge, the app bridge — for the three doors in Sections 2 and 3, in the deck's left-to-right order. The deck and the video are now telling the same three-channel story, and matching the icons is the cheapest way to make a judge feel they already know the system when the slide comes up.

Citation cards are mandatory. Every statistic gets its source in small type underneath, on screen, for as long as the number is up. This is the cheapest credibility in the whole edit and most teams skip it.

B.5 — Subtitles

Burn in **English subtitles for the entire video**, not just the Hindi. Judges may watch on a laptop in a noisy hall. Hindi lines get Devanagari plus an English gloss on a second line.

PART C — RECOMMENDATIONS

1. **Film the IVR call first, before anything else is written or shot.** It is the only real asset in the video and everything else is scheduled around it. If it looks good, Section 3 can lean on it harder. If it doesn't, we need to know while there's still time to fix it.
 2. **Section 4 is the differentiator — protect it.** Every team will show a voice assistant. Almost none will name the DL-PACC, the Perspective Plan, and four statutory dates. If you need to find time, take it from Section 2, not Section 4.
 3. **Section 5's dialect admission is a feature, not a confession.** Brief the whole team on this, because the instinct under pressure is to soften it. Do not soften it. It is the most memorable thirty seconds we have.
 4. **Cast Section 4 and Section 1 first.** Section 4 needs the most confident speaker. Section 1 needs whoever can deliver "Safai Karmchari" flat without flinching. Sections 2, 3 and 5 can be anyone.
 5. **Record all five separately, same wall, same lens, same height.** Cutting between five rooms is the fastest way to make a student video look like a student video.
 6. **One safety take per person, 20% slower.** The 60-second sections always run long on the first pass, and pacing cannot be fixed in the edit without cutting content.
 7. **Brief everyone on the five ⚠ traps before the shoot, not during:** no field-work claims, no dialect overclaim, no PM-AJAY portal integration claim, no "district submits the Perspective Plan", and no "works offline" said about the app's microphone.
 8. **Record the WhatsApp and app screens before the shoot day, not after.** They are short beats, which makes them easy to postpone and easy to end up faking under time pressure. Both builds run today; record them while that is cheap. If the WhatsApp voice-note decode is still open on shoot day, record the exchange anyway — the message goes, the audio reply comes back, and that is what is on screen.
-

PART D — Q&A BANK

Not in the video. Prepare these; each has been raised or is likely to be.

"MSDE already shipped a multilingual WhatsApp skilling assistant with Meta and Sarvam. How is this different?"

That one answers questions. Ours conducts a structured interview against seven mandated fields, gates on published NSQF entry requirements, and emits a district-level planning artefact on the scheme's statutory calendar. It is the officer half that makes it a PM-AJAY system rather than a chatbot. — *Rehearse this as one paragraph. Do not improvise it.*

"Why is PMKVY evidence relevant to PM-AJAY?"

PM-AJAY's GIA skilling is executed through the same SSDMs, Sector Skill Councils, NSQF qualification packs and MSDE Common Norms. The failure modes are inherited.

"The PM-AJAY portal already has district and state dashboards. Why build yours?"

We don't replace it. The portal has no way to collect a non-literate beneficiary's aspirations by voice — that is the gap, and that is our half. We produce what the officer uploads.

"How many languages do you support?"

Languages, several. Dialects — models do exist now, including on Bhashini, and they are still wrong around three words in ten. We don't claim to have improved on that. The system is built so that a wrong transcript still produces a right field, and we publish the error rate alongside the extraction accuracy rather than quote one without the other.

"Which of the three channels have you actually built?" (*Ask this of yourselves before a judge does.*)

All three run. The phone channel is measured over 1,822 turns. WhatsApp goes direct to Meta's Cloud API — HMAC-verified webhook, one question per message, replies as audio addressed by media id — and the voice-note decode is the last step, which is on our published FUTURE list. The app is a Capacitor Android build with the same interview and an offline store. — *Answer in that order. Naming the one open step yourself is what makes the other two claims land.*

"Does the app work offline?"

The interview and the record do — answers are held on the device and sync from an outbox when the link comes back. Speech does not yet: it goes through Android's own recogniser, which usually needs a network hop. The on-device path is Vosk behind a Capacitor plugin, and it is the single remaining native piece. — *Never answer this one with a flat yes.*

"Why WhatsApp as well as a call? Isn't that the same thing twice?"

Different failure modes. The call is synchronous and needs her free for four minutes; WhatsApp is asynchronous and she answers between chores. And one is a feature phone, the other is not. Same engine, same seven fields, same eligibility gate behind both.

"Who holds the phone if rural women don't own one?"

The kiosk channel, operated by the CSC Village Level Entrepreneur. It is also the problem statement's own answer to "inadequate technical and support team at ground level."

"Does experience really substitute for schooling under NSQF?" (*Most likely question from an NCVET-aware juror. Know the limits.*)

It trades against it, it doesn't replace it. The gazette's Level 2.5 entry row admits a 9th-grade pass with no experience, or a 5th-grade pass with four years of relevant work — same door. But above Level 2 there is always a minimum grade alongside the years, the experience requirement caps at five years, and **beyond Level 2.5 a formal experience certificate is explicitly still required**. The genuinely qualification-free route is RPL, and that needs an NCVET-guideline assessment. Our engine encodes the published table; it does not invent routes through it.

"What's your accuracy?"

Two numbers, always given together: word error rate, and field-extraction accuracy. They diverge, and the divergence is the design. — *Have the real figures on a card before the pitch. Do not quote a number you have not measured.*